

A Course End Project

on

**OBJECT
REPLICATION
WITH AWS**

Submitted in the Partial Fulfillment of
the Requirements

for the Award of the Degree of

BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING (AI&ML)

Submitted

By

Team No.: 11

B.Raju **21881A6673**

Saisrikar Devasani **21881A6679**

R. Sai Siddarth **21881A66B5**

Under the Esteemed Guidance of

Dr. P. Pavan Kumar

Associate Professor



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING (AI&ML)

VARDHAMAN COLLEGE OF ENGINEERING
(AUTONOMOUS)

Affiliated to **JNTUH**, Approved by **AICTE**, Accredited by **NAAC**, with **A++** Grade, **ISO 9001:2015** Certified
Kacharam, Shamshabad, Hyderabad – 501218, Telangana, India

2023- 24

ACKNOWLEDGEMENT

The satisfaction that accompanies the successful completion of the task would be put incomplete without the mention of the people who made it possible, whose constant guidance and encouragement crown all the efforts with success.

We wish to express my deep sense of gratitude to **Dr. P. Pavan Kumar, Associate Professor** for their able guidance and useful suggestions, which helped us in completing the design part of potential project in time.

We are particularly thankful to **Dr. M. A. Jabbar**, Professor & Head, Department of Computer Science and Engineering (AI&ML) for his guidance, intense support and encouragement, which helped us to mould our project into a successful one.

We show gratitude to our honourable Principal **Dr. J. V. R. Ravindra**, for having provided all the facilities and support.

We avail this opportunity to express our deep sense of gratitude and heartfelt thanks to **Dr. Teegala Vijender Reddy**, Chairman and **Sri Teegala Upender Reddy**, Secretary of VCE, for providing congenial atmosphere to complete this project successfully.

We also thank all the staff members of the department of **CSE(AI&ML)** for their valuable support and generous advice. Finally, thanks to all our friends and family members for their continuous support and enthusiastic help.

B. Raju 21881A6673

Saisrikar Devasani 21881A6679

R. Sai Siddarth 21881A66B5



VARDHAMAN COLLEGE OF ENGINEERING

(AUTONOMOUS)

Affiliated to **JNTUH**, Approved by **AICTE**, Accredited by **NAAC**, with **A++** Grade, **ISO 9001:2015** Certified
Kacharam, Shamshabad, Hyderabad – 501218, Telangana, India

Department of Computer Science and Engineering (AI&ML)

CERTIFICATE

This is to certify that the Course End Project report work entitled “**Object Replication With AWS**” carried out by Mr. **B. Raju**, Roll Number **21881A6673**, Mrs. **Saisrikar Devasani**, Roll Number **21881A6679**, Mr. **R. Sai Siddarth**, Roll Number **21881A66B5** towards Course End Project and submitted to the Department of Computer Science and Engineering(AI&ML), in partial fulfillment of the requirements for the award of degree of **Bachelor of Technology** in **Computer Science and Engineering (AI&ML)** during the year 2023-24.

Name & Signature of the Instructors

Dr. P. Pavan Kumar

Associate Professor

Name & Signature of the HOD

Dr. M A Jabbar

HOD, CSE(AI&ML)

ABSTRACT

Object replication with AWS (Amazon Web Services) S3 offers organizations a robust solution for enhancing data availability, durability, and compliance across distributed environments. This approach involves replicating data objects stored in Amazon S3 buckets across multiple AWS regions, ensuring resilience against regional failures and reducing latency for global users. Key considerations include configuring cross-region replication policies, implementing security measures such as encryption and access controls, and optimizing costs through lifecycle management of replicated data. Object replication with AWS supports compliance with data residency regulations and strengthens disaster recovery capabilities, making it essential for modern cloud architectures seeking to balance performance, resilience, and regulatory requirements.

Keywords: AWS S3, object replication, cross-region replication, data durability, data availability, compliance, disaster recovery, AWS regions, data sovereignty, cost optimization.

ABBREVIATIONS

| Abbreviation | Expansion |
|---------------------|----------------------------|
| VM | Virtual Machine |
| AWS | Amazon Web Services |
| ACL | Access control list |
| VHD | Virtual Hard Disk |

List of Figures

| Figure Name | Title |
|--------------------|-----------------------------------|
| Figure 1-3 | Create Buckets |
| Figure 4-7 | Configure Buckets |
| Figure 8-12 | Set Up Replication |
| Figure 13 | Upload Object |
| Figure 14 | Verification Of the Object |

Table of Contents

| Chapt er No. | Title | Pag e No. |
|-----------------------------|--|--------------------------|
| | Acknowledgement | 2 |
| | Abstract | 4 |
| | Abbreviations | 5 |
| | List of Figures | 5 |
| Chapter 1 | Introduction | 7-8 |
| | 1.1 Background | 7 |
| | 1.2 Importance of Object replication | 7 |
| | 1.3 Scope of The project | 7 |
| | 1.4 Primary Objectives | 7 |
| | 1.5 Secondary Objectives | 7 |
| | Object Replication with AWS | 8 |
| Chapter 2 | Problem Definition and Proposed System Methodology | 9-16 |
| | 2.1 Problem Statement | 9 |
| | 2.2 Requirements for Effective Replication | 9 |
| | 2.3 Steps | 10-16 |
| Chapter 3 | Software Requirements Specification and Hardware Requirements | 17 |
| | 3.1 Software Requirements Specification | 17 |
| | 3.2 Hardware Requirements | 17 |
| | 3.3 Browser Requirements | 17 |
| Chapter 4 | Results and Discussions | 18- 19 |
| Chapter 5 | Conclusion and Future Scope | 20 |
| | References | 21 |

Chapter 1

Introduction

1.1 Background

In today's digital era, data availability, durability, and accessibility are crucial for businesses and organizations. Object replication is a key feature in cloud storage solutions, allowing for automatic copying of data across different geographical locations. AWS (Amazon Web Services) provides a robust object replication service through Amazon S3 (Simple Storage Service), ensuring that data is replicated to improve redundancy, minimize latency, and facilitate disaster recovery.

1.2 Importance of Object Replication

Object replication in AWS S3 ensures data is consistently copied to different regions or within the same region, enhancing data durability and availability. It helps in achieving compliance with data residency requirements, improves access times for distributed applications, and provides a reliable backup in case of accidental data loss or regional failures.

1.3 Scope of the Report

This report explores the concept of object replication within AWS, focusing on the functionalities, benefits, and implementation strategies of AWS S3 replication. It covers the setup process, best practices, potential challenges, and performance analysis to provide a comprehensive understanding of AWS object replication.

Objectives

1.4 Primary Objectives

- To understand the features and capabilities of AWS S3 replication.
- To implement a replication strategy using AWS S3.
- To evaluate the performance and cost implications of using AWS S3 replication.

1.5 Secondary Objectives

- To compare AWS S3 replication with other cloud providers.
- To identify best practices for optimizing replication performance.
- To explore real-world use cases and scenarios for AWS S3 replication.

Object Replication with AWS

Object replication in AWS involves creating copies of data objects, such as files or volumes, across different AWS Regions or Availability Zones to ensure data redundancy, improve availability, and support disaster recovery strategies. This process allows businesses to replicate critical data closer to users or in different geographic locations for compliance or performance reasons.

Key Features

- 1. Cross-Region Replication:** AWS enables replication of objects across different AWS Regions, allowing businesses to maintain copies of their data in multiple geographic locations for disaster recovery and compliance.
- 2. Data Availability:** By replicating objects across AWS Availability Zones within a Region, AWS ensures high availability and fault tolerance. Availability Zones are physically separated data centers within a Region.
- 3. Automated Replication:** AWS offers tools and features to automate the replication process, reducing manual intervention and ensuring consistency and efficiency in data replication tasks.
- 4. Data Integrity and Security:** AWS provides robust security measures, including encryption in transit and at rest, IAM (Identity and Access Management) policies, and audit logging, to ensure data integrity and protect against unauthorized access.
- 5. Performance Optimization:** Replicating data closer to end-users reduces latency and improves application performance, enhancing user experience for globally distributed applications.
- 6. Cost Efficiency:** AWS offers flexible pricing models for data transfer and storage, allowing businesses to optimize costs associated with object replication based on their specific needs and usage patterns.

Object replication with AWS provides businesses with the flexibility and scalability needed to ensure data resilience, compliance, and optimal performance across diverse geographic regions.

Chapter 2

Problem Statement and Proposed System Methodology

1. Problem Statement

Data replication across multiple locations presents several challenges, including network latency, data consistency, security concerns, and compliance with regional data laws. Ensuring seamless and efficient replication without impacting application performance is a critical issue for many organizations.

2. Requirement for Effective Replication.

Effective data replication requires a secure and scalable infrastructure, minimal latency, robust monitoring and logging capabilities, and compliance with regulatory standards. AWS S3 provides the necessary tools and features to meet these requirements, making it a preferred choice for many businesses.

3. Steps

Here are the detailed steps for creating and verifying a snapshot in Azure:

Step 1: Create Buckets

- In the AWS Management Console, navigate to Amazon S3 and create two buckets in the same AWS region. Ensure both buckets are configured with similar settings.

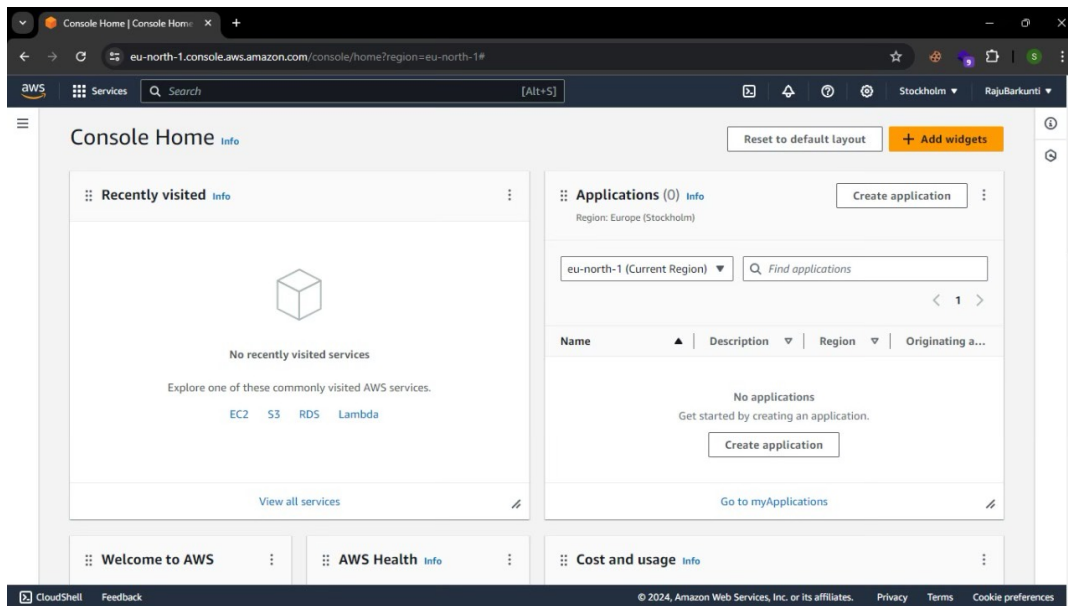
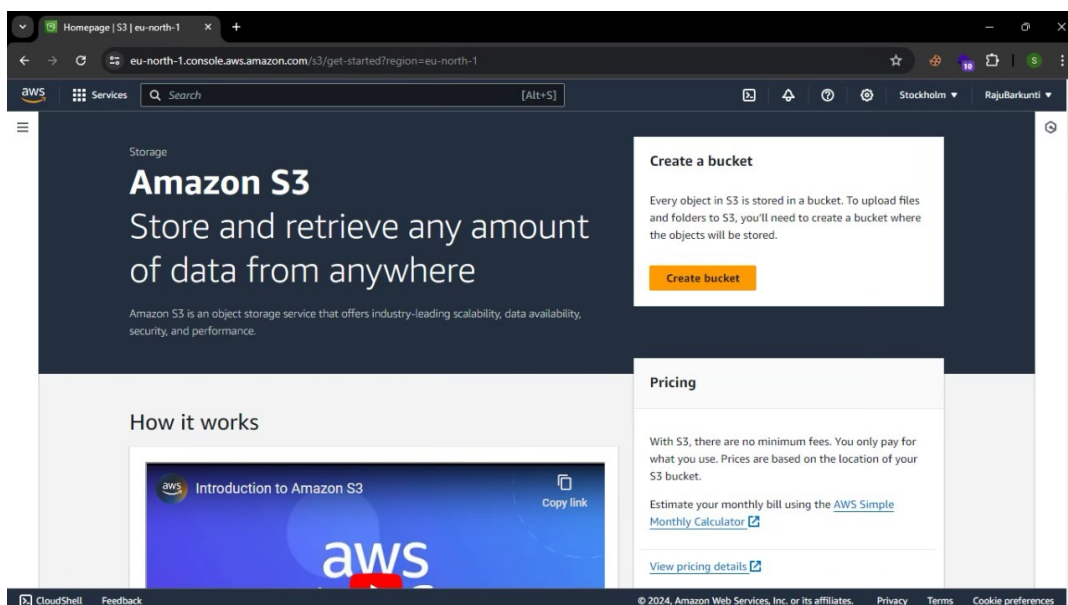


Figure:1 & 2



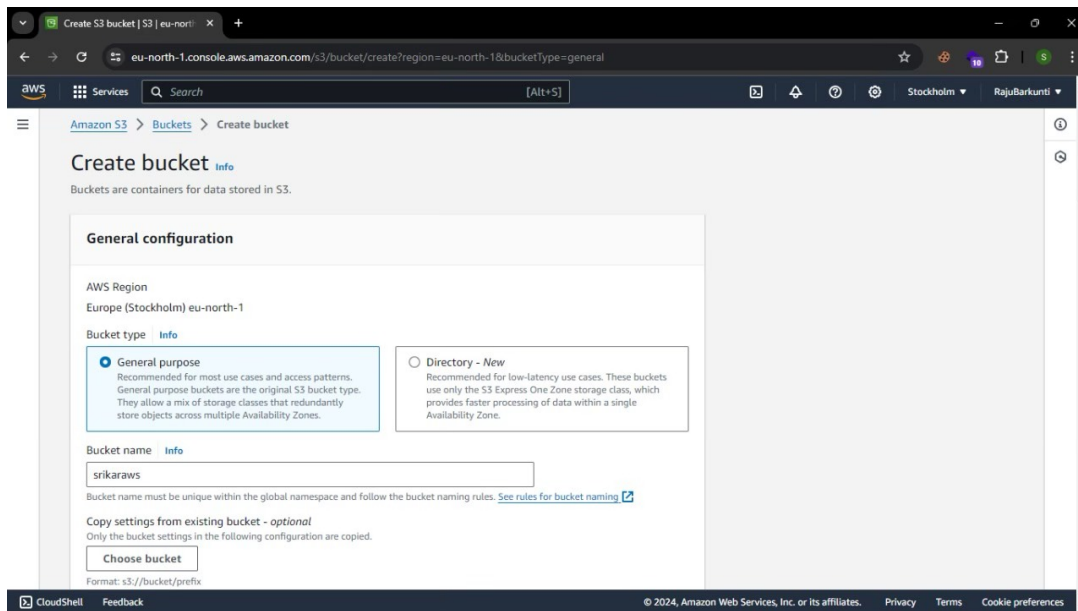


Figure: 3

Step 2: Configure Buckets

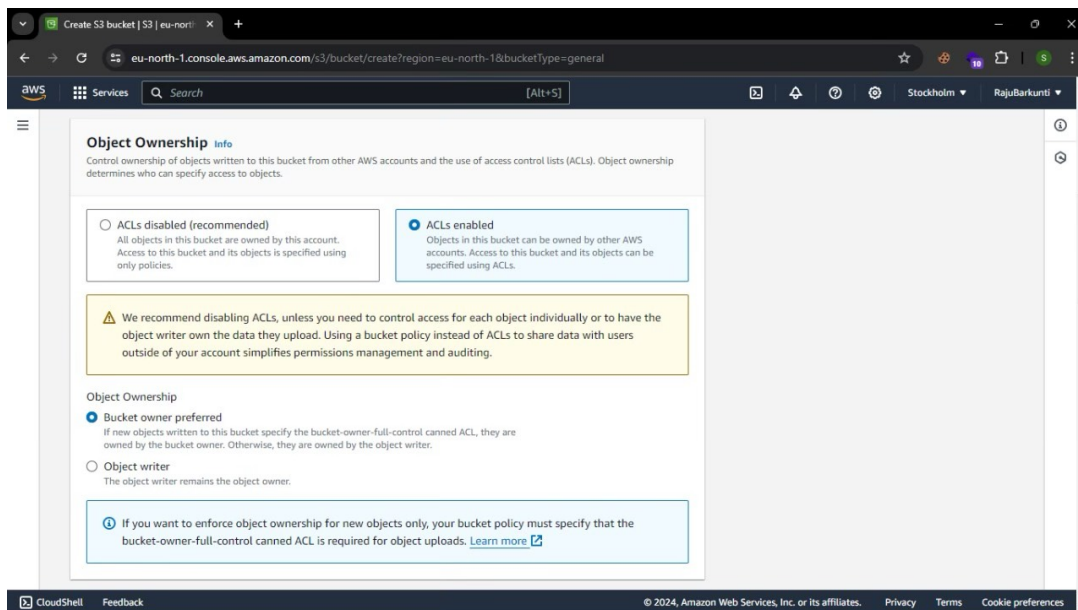


Figure:4

- Enable ACLs: Select "ACLs Enabled" for both buckets to manage permissions using AccessControl Lists (ACLs).

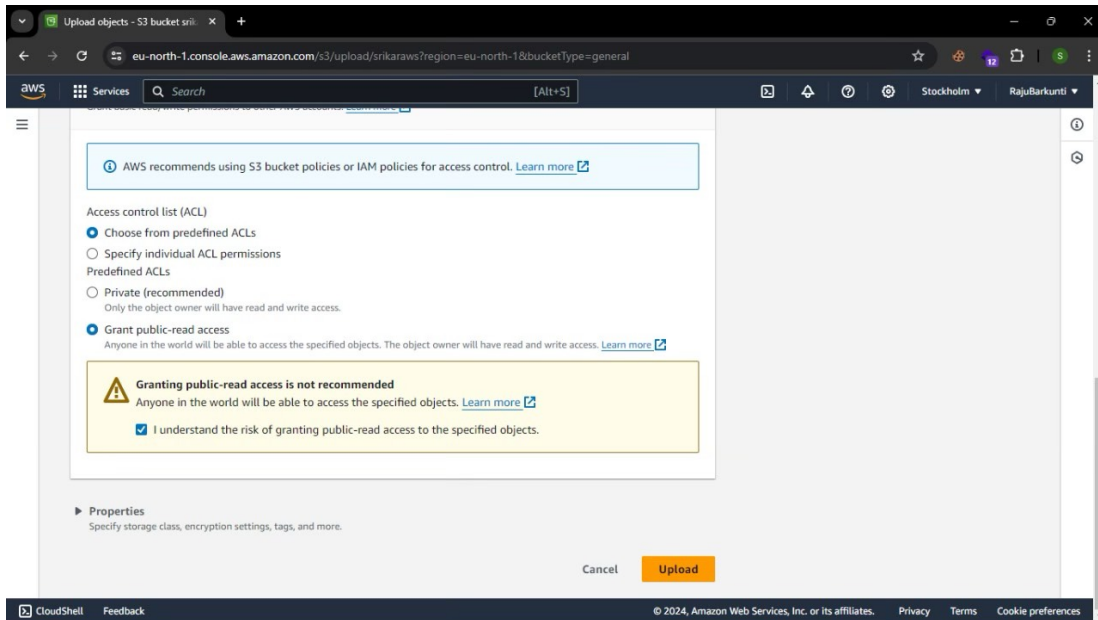


Figure:5

- Set Public Access: Configure the buckets to allow public access as required by your use case.

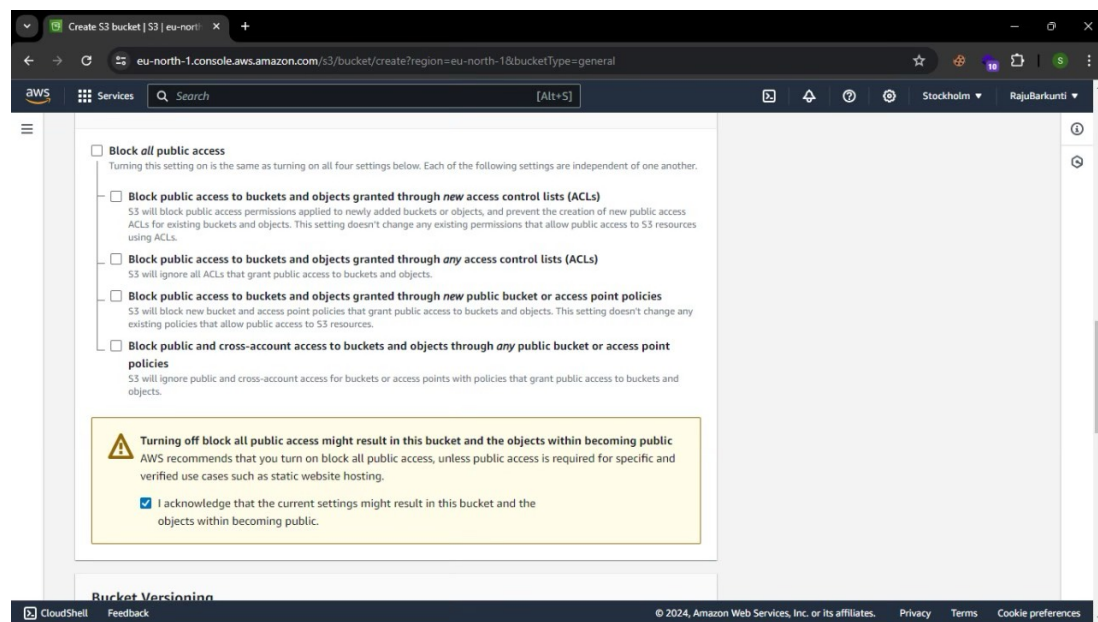


Figure:6

- Enable Versioning: Go to the "Properties" tab of each bucket, click "Edit" for "Bucket Versioning," select "Enable," and save the changes. Versioning is essential for replication as it ensures that every version of an object is preserved.

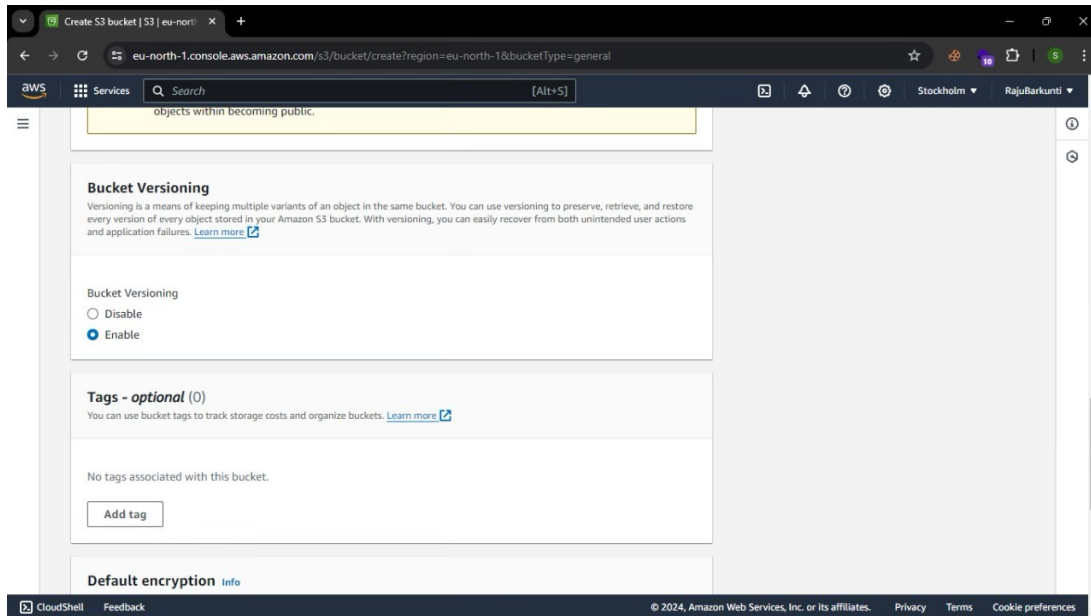


Figure:7

Step 3: Set up Replication

- Go to the Source Bucket: Click on the source bucket, navigate to the "Management" tab, and select "Replication Rules."
- Create Replication Rule:
 - - Click "Create replication rule," enter a name for the rule, and set the status to "Enabled."
 - - Specify that the rule applies to all objects in the bucket.

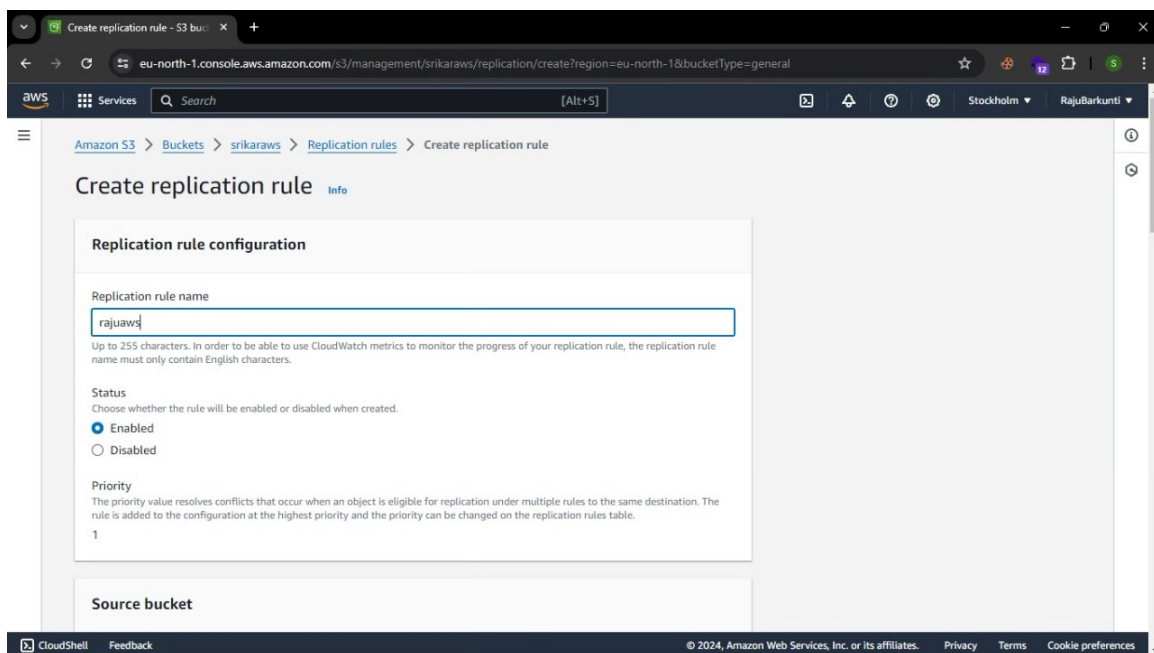


Figure:8

- Select Destination Bucket:
 - - Choose the destination bucket for replication.
 - - Select "Create new role" to allow AWS to manage the necessary permissions.

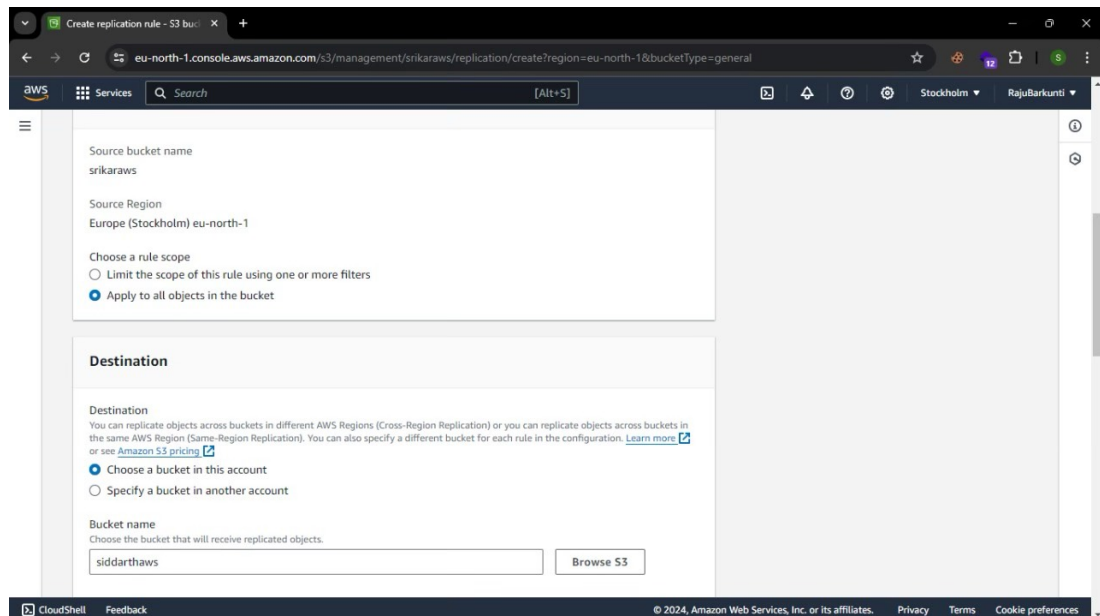


Figure:9

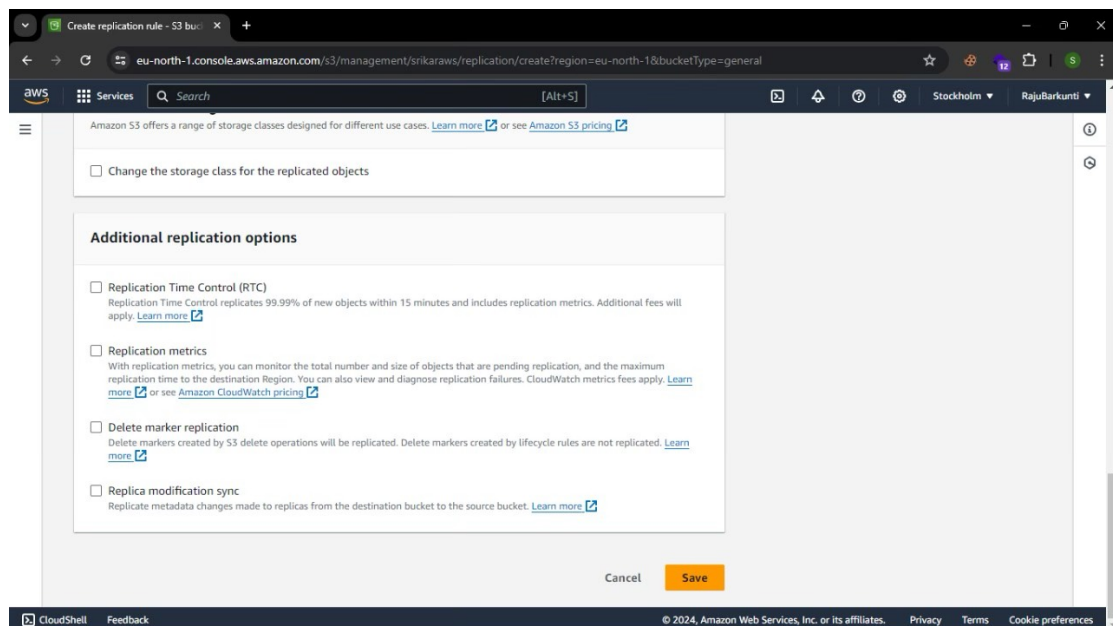


Figure:10

- Save and Apply
 - - Save the rule. Choose "No, do not replicate existing objects" if you only want to replicate new objects.
 - - Click "Submit" to finalize the replication rule.

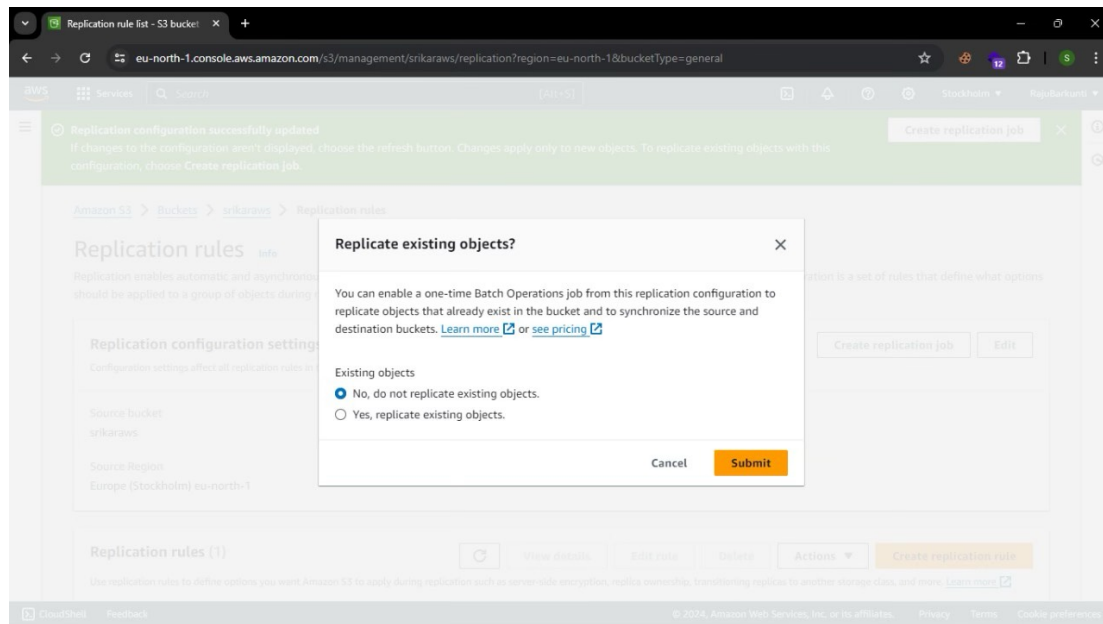


Figure:11

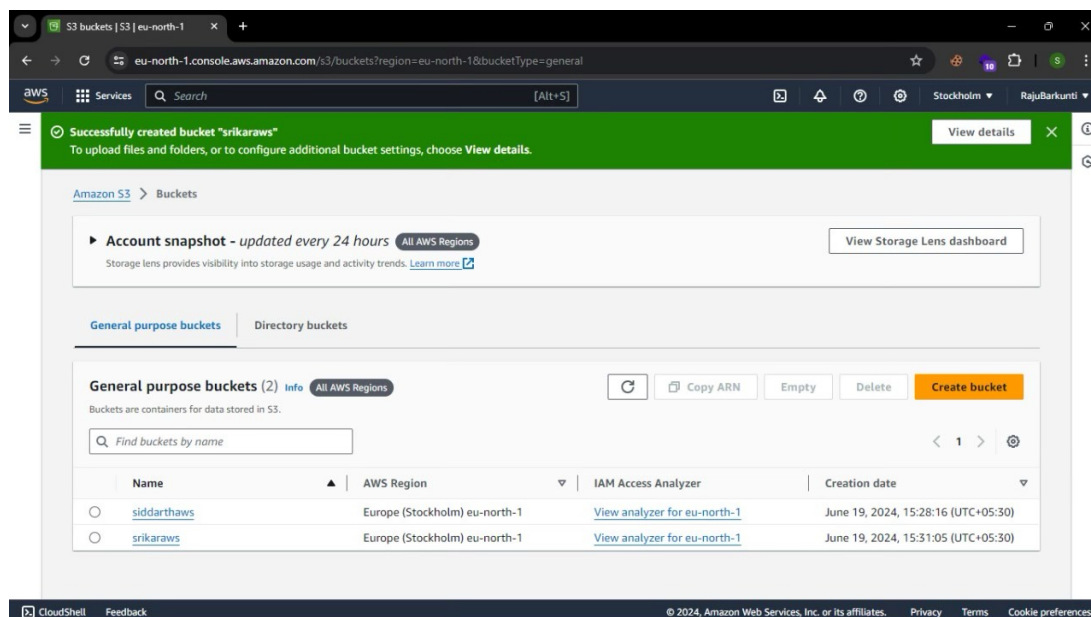


Figure:12

Step 4: Verify Replication

- Upload Objects:
 - - Upload objects to the source bucket and grant public permissions as required

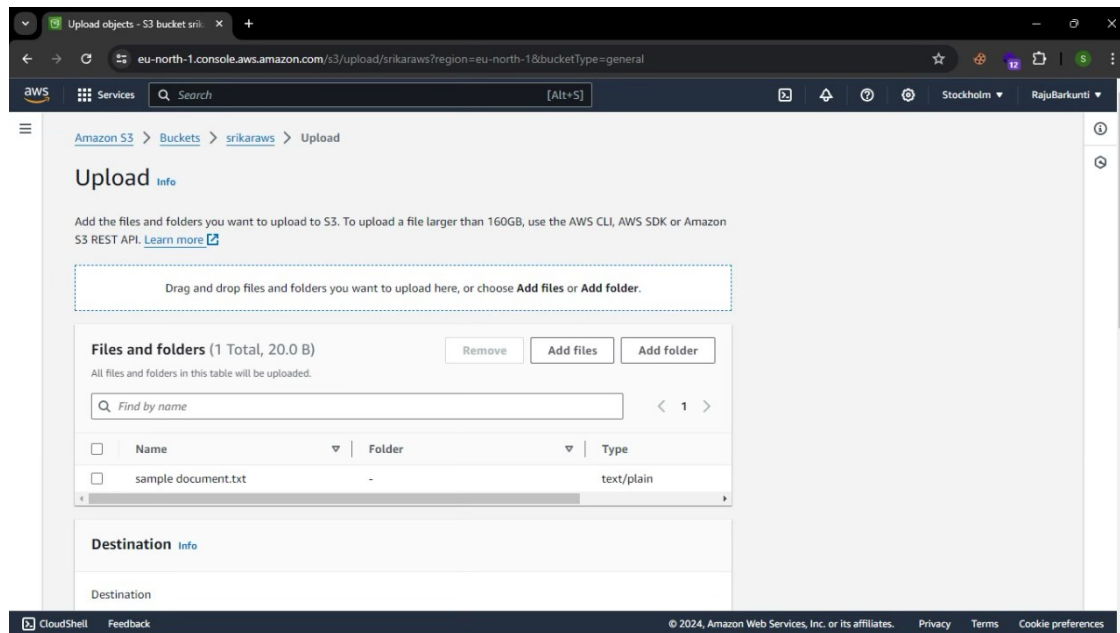


Figure:13

- Check Replication:
 - - Verify that the objects are replicated to the destination bucket by checking the destination bucket's contents.

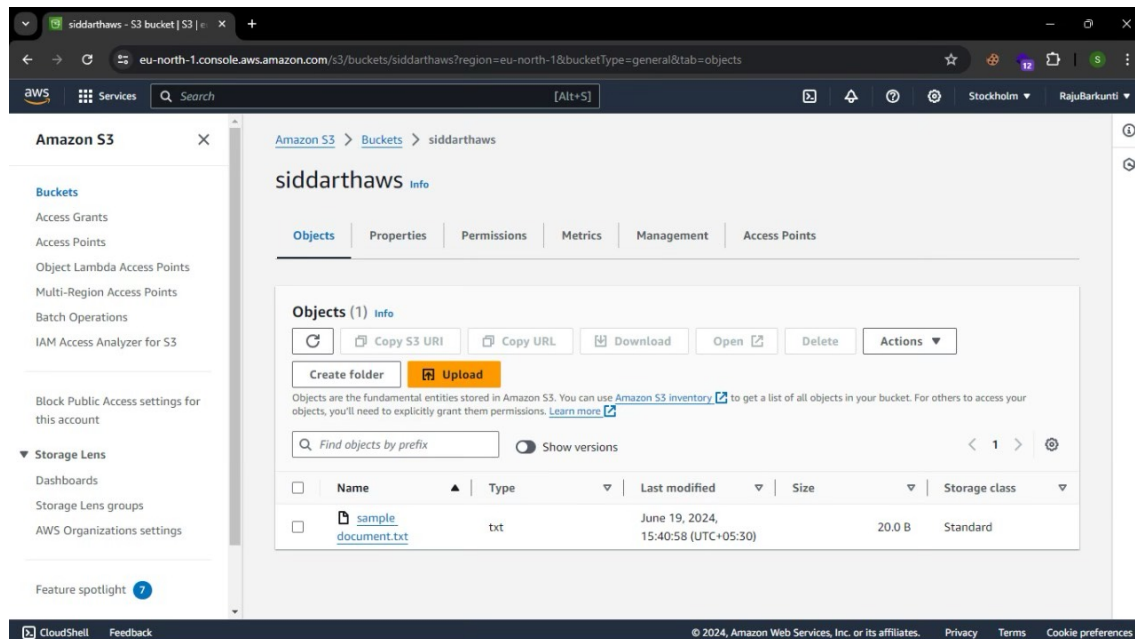


Figure:14

Chapter 3

Software Requirements Specification and Hardware Requirements

3.1 Software Requirements:

- **AWS Account:** You need an active AWS account with appropriate permissions to create and manage resources.
- **AWS Management Console or AWS CLI:** You can use the AWS Management Console web interface or AWS Command Line Interface (CLI) for setting up and managing services.
- **S3 Buckets:** Source and destination S3 buckets configured with versioning enabled.
- **AWS IAM Roles and Policies:** IAM roles and policies to grant necessary permissions for replication.
- **Network Configuration:** Ensure proper network connectivity and configurations for S3 bucket access and cross-region replication.

3.2 Hardware Requirements:

- **Client Device:** Any computer or device capable of running a modern web browser for accessing the AWS Management Console.
- **Internet Connectivity:** Stable internet connection for accessing AWS services and managing configurations.
- **Storage Requirements:** Depending on the volume of data to be replicated, sufficient storage space in S3 buckets and possibly on local devices for managing data transfer and caching.

Additional Considerations:

- **AWS Region Selection:** Choose appropriate AWS regions for your source and destination buckets based on latency, compliance, and cost considerations.
- **Data Transfer Costs:** Be aware of potential data transfer costs associated with cross-region replication in AWS.

3.3 Browser Requirements:

- **Web Browser:** Use of a compatible web browser (such as Google Chrome, Mozilla Firefox, Microsoft Edge, or Safari) to access and navigate through the Azure Portal.

These hardware requirements ensure that you can effectively access and manage Azure resources through the Azure Portal, perform tasks like creating and managing snapshots, uploading data to storage accounts, and deploying virtual machines in the desired Azure region (Southeast Asia in this case). The actual computation and processing are handled by Microsoft Azure's cloud infrastructure, making the hardware requirements minimal for the user's local device.

Chapter 4

Results and Discussions

Results:

Improved Data Redundancy and Durability:

Object replication with AWS S3 helps in achieving improved redundancy by storing copies of objects across different AWS regions. This ensures data durability and availability even in the event of regional failures.

Reduced Latency for Global Users:

By replicating objects closer to users in different geographic regions, you can reduce latency for accessing objects. This improves user experience and application performance, especially for global user bases.

Compliance and Data Residency Requirements:

Object replication allows you to comply with data residency requirements by storing copies of data in specific regions or countries where data sovereignty laws apply. This helps in meeting regulatory and compliance standards.

Cost Efficiency:

AWS S3 offers cost-effective storage solutions for object replication. You can optimize costs by leveraging S3's lifecycle policies to transition objects to lower-cost storage classes after replication.

Discussions:

Data Consistency and Synchronization:

Discuss the mechanisms AWS provides to ensure data consistency and synchronization between source and destination buckets during object replication. This includes understanding eventual consistency models and how replication configurations impact data integrity.

Security and Access Control:

Explore AWS IAM roles and policies used to secure access to replicated objects. Discuss best practices for securing S3 buckets, enabling encryption, and managing access permissions across different regions.

Performance Considerations:

Evaluate the performance implications of object replication, such as data transfer speeds, network latency, and AWS Direct Connect options for optimizing replication performance. Discuss strategies for monitoring and optimizing replication workflows.

Operational Considerations and Automation:

Consider operational aspects such as monitoring replication status, handling replication failures, and implementing automated workflows using AWS Lambda functions or AWS CloudFormation templates. Discuss strategies for managing replication configurations at scale.

Business Continuity and Disaster Recovery:

Object replication plays a crucial role in business continuity and disaster recovery planning. Discuss how AWS S3 object replication contributes to maintaining data availability and disaster recovery readiness in multi-region architectures.

Chapter 5

Conclusion and Future Scope

Conclusions:

Object replication with AWS S3 offers significant benefits in terms of enhancing data availability, durability, and compliance across global infrastructures. By replicating objects across multiple AWS regions, organizations can ensure high resilience against regional failures and improve data access speeds for users worldwide. This approach not only strengthens disaster recovery capabilities but also supports compliance with diverse data sovereignty regulations by enabling data storage in specific geographic locations.

Furthermore, the cost-efficiency of AWS S3's storage options, coupled with lifecycle management policies, allows organizations to optimize expenses while maintaining robust data redundancy through effective management of replicated data. This combination of enhanced data resilience, improved performance, compliance adherence, and cost optimization makes AWS S3 object replication a compelling solution for modern cloud architectures.

Future Scope:

Looking ahead, there are several avenues for further enhancing object replication strategies with AWS S3:

- **Advanced Automation:** Expanding automation capabilities using AWS Lambda and CloudFormation to streamline replication workflows, monitor replication status in real-time, and dynamically adjust configurations based on workload demands.
- **Security Enhancements:** Implementing more robust security measures such as encryption, fine-grained access controls, and comprehensive auditing to fortify the protection of replicated data across regions.
- **Integration with AI/ML:** Exploring integration opportunities with AI and ML technologies to leverage replicated data for predictive analytics, automated decision-making processes, and enhanced business insights.
- **Real-time Replication:** Investigating solutions for achieving near real-time object replication to minimize recovery point objectives (RPOs) and ensure synchronized data availability across distributed environments.
- **Hybrid Cloud Integration:** Developing strategies to seamlessly integrate AWS S3 object replication with hybrid cloud architectures, including on-premises data centers and other cloud providers, to establish unified data management frameworks.

References

1. **AWS Documentation - Cross-Region Replication (CRR):** This official AWS documentation provides comprehensive information on setting up and managing cross-region replication for Amazon S3. [AWS Documentation - Cross-Region Replication](#)
2. **AWS Blog - Best Practices for Amazon S3 Cross-Region Replication:** This AWS blog post discusses best practices and considerations for implementing cross-region replication with Amazon S3. [AWS Blog - Best Practices for Amazon S3 Cross-Region Replication](#)
3. **AWS Online Tech Talks - Introduction to Amazon S3 Replication:** This AWS online tech talk provides an introduction to Amazon S3 replication features and how to use them effectively. [AWS Online Tech Talks - Introduction to Amazon S3 Replication](#)
4. **AWS Whitepaper - Architecting for Cross-Region Replication in Amazon S3:** This whitepaper from AWS discusses architectural considerations and best practices for designing cross-region replication solutions using Amazon S3. [AWS Whitepaper - Architecting for Cross-Region Replication in Amazon S3](#)
5. **AWS re**
2020 - Deep Dive into Amazon S3 Replication: This session from AWS re
2020 provides a deep dive into Amazon S3 replication, covering advanced topics and use cases. [AWS re](#)
[2020 - Deep Dive into Amazon S3 Replication](#)

GITHUB LINK: <https://github.com/SaiSiddarth21/CCV>